

Feedforward Neural Networks and Backpropagation

Changyou Chen

Department of Computer Science and Engineering
University at Buffalo, SUNY
`changyou@buffalo.edu`

February 21, 2019

Backpropagation

Importance of Backpropagation

- ➊ Backprop is a technique for computing derivatives quickly by avoiding duplicated computations.
- ➋ It is the key algorithm that makes training deep models computationally tractable.
- ➌ For modern neural networks it can make training gradient descent 10 million times faster relative to naive implementation:
 - ▶ it is the difference between a model that takes a week to train instead of 200,000 years.
- ➍ It is based on the computational graph.
- ➎ Simply called backprop.

Anecdote


Geoffery Hinton and BP

- **1986:** Invented Backpropagation.
- **2017:** “Discard it, do it from the beginning!”
 - Capsule networks.



Basic Steps in Backpropagation

- 1 Two passes through the computational graph:
 - ▶ forward pass: compute the outputs of each layer.
 - ▶ backward pass: compute the gradients of parameters in each layer, based on results from the forward pass.
- 2 Backprop algorithm does this using a simple and inexpensive procedure.



Backprop vs. Learning


- 1 Backprop does not mean the whole learning algorithm.
- 2 Backprop only refers to the method of computing the gradient.
 - ▶ Another algorithm, such as SGD, is used to perform learning using this gradient:

$$\mathbf{w}_t = \mathbf{w}_{t-1} - \eta \nabla_{\mathbf{w}} J(\mathbf{w})$$

- 3 Often misunderstood to being specific to multilayer neural networks.
 - ▶ It can be used to compute derivatives for any function.
 - ▶ It also can be used to compute Jacobian of a function with multiple outputs.
 - ▶ We restrict to case where f has a single output.

Recap: Computational Graph


- Each node is either
 - ▶ a variable: scalar, vector, matrix, tensor, or other type
 - ★ simple function of one or more variables
 - ★ functions more complex than operations are obtained by composing operations
 - ▶ or an operation
 - ★ if variable y is computed by applying operation to variable x then draw directed edge from x to y .



Example: Graph of Logistic Regression

$$\hat{y} = \sigma(\mathbf{w}^T \mathbf{x} + b)$$


- Variables in graph $\mathbf{u}^{(1)}$ and $\mathbf{u}^{(2)}$ are not in the original expression, but are needed in the graph.



Example: Graph of ReLU

$$\mathbf{H} = \max \left\{ 0, \mathbf{W}^T \mathbf{X} + \mathbf{b} \right\}$$

- Variables in graph $\mathbf{U}^{(1)}$ and $\mathbf{U}^{(2)}$ are not in the original expression, but are needed in the graph.



Recap: Calculus' Chain Rule for Scalars

- Let x be a real number; f and g be functions mapping from a real number to a real number.
- Let $y = g(x)$, $z = f(g(x)) = f(y)$, then

$$\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$$

Generalizing Chain Rule to Vectors

- Let $\mathbf{x} \in \mathbb{R}^m$, $\mathbf{y} \in \mathbb{R}^n$; $f : \mathbb{R}^m \rightarrow \mathbb{R}^n$, and $g : \mathbb{R}^n \rightarrow \mathbb{R}$.
- Let $\mathbf{y} = g(\mathbf{x})$, $z = f(g(\mathbf{x})) = f(\mathbf{y})$, then

$$\frac{dz}{dx_i} = \sum_j \frac{dz}{dy_j} \cdot \frac{dy_j}{dx_i}$$

- ▶ Like multiple paths when each node represents one element of a vector.
- In vector notation, this is

$$\nabla_{\mathbf{x}} z = \left(\frac{\partial \mathbf{y}}{\partial \mathbf{x}} \right)^T \nabla_{\mathbf{y}} z$$

where $\frac{\partial \mathbf{y}}{\partial \mathbf{x}}$ is the $n \times m$ Jacobian matrix of g .

- Backprop algorithm consists of performing “Jacobian-gradient” product for each step of graph.

Generalizing Chain Rule to Tensors

- 1 Backpropagation is usually applied to tensors with arbitrary dimensionality.
- 2 This is almost the same as with vectors:
 - ▶ we could flatten each tensor into a vector, compute a vector-valued gradient and reshape it back to a tensor.
- 3 In this view backpropagation is still multiplying Jacobians by gradients:

$$\nabla_{\mathbf{x}} z = \left(\frac{\partial \mathbf{y}}{\partial \mathbf{x}} \right)^T \nabla_{\mathbf{y}} z$$

Chain Rule for Tensors

- 1 To denote gradient of value z w.r.t. a tensor \mathbf{X} , we write as if \mathbf{X} were a vector:
 - ▶ an element in a k -order tensor is indexed by a tuple (i_1, \dots, i_k) , we abstract this away by a single (vector) variable $i \triangleq (i_1, \dots, i_k)$ to represent the tuple.
- 2 For all possible tuples i , $(\nabla_{\mathbf{X}} z)_i$ gives $\frac{\partial z}{\partial X_i}$:
 - ▶ exactly same as how for all possible indices i in a vector, $(\nabla_{\mathbf{X}} z)_i$ gives $\frac{\partial z}{\partial X_i}$.
- 3 Chain rule for tensors: let $\mathbf{Y} = g(\mathbf{X})$ and $z = f(\mathbf{Y})$:

$$\nabla_{\mathbf{X}} z = \sum_i \underbrace{(\nabla_{\mathbf{X}} Y_i)}_{\text{tensor}} \underbrace{\frac{\partial z}{\partial Y_i}}_{\text{scalar}}$$

Backprop is Recursive Chain Rule

- 1 Backprop is obtained by recursively applying the chain rule.
- 2 Using the chain rule, it is straightforward to write expression for gradient of a scalar w.r.t. any node in graph producing that scalar.
- 3 However, evaluating that expression on a computer has some extra considerations:
 - ▶ e.g., many subexpressions may be repeated several times within overall expression.
 - ▶ should we store subexpressions or recompute them?


Example of repeated subexpressions

- ❶ Let ω be the input to the graph; Use same function $f : \mathbb{R} \rightarrow \mathbb{R}$ at every step: $x = f(\omega)$, $y = f(x)$, $z = f(y)$.

$$\frac{\partial z}{\partial \omega} = \frac{\partial z}{\partial y} \cdot \frac{\partial y}{\partial x} \cdot \frac{\partial x}{\partial \omega} = f'(\textcolor{red}{y})f'(\textcolor{blue}{x})f'(\omega) \quad (1)$$


$$= f'(\textcolor{red}{f(f(\omega))})f'(\textcolor{blue}{f(\omega)})f'(\omega) \quad (2)$$

- ❷ We can use (1) to compute the derivative, where $f(\omega)$ is computed once and stored it in x .
- ▶ This is the approach taken by backprop.
- ❸ We can also use (2) to compute the derivative, where $f(\omega)$ is recomputed each time it is needed.
- ▶ For low memory, (1) preferable: reduced runtime.
 - ▶ (2) is also valid chain rule, useful for limited memory.
- ❹ For complicated graphs as in deep neural networks, (2) exponentially wastes computations for repeated subexpressions.



Basic Backprop Algorithm

- Assume each u_i is a scalar.
- Each $\{u_j\}$ is connected by a subset of other $\{u_k\}_{k \leq j}$, denoted as $Pa(u_j)$.
- For Each $\{u_j\}_{j > i}$, $u_j = f^{(j)}(Pa(u_j))$.
- The basic backprop consists of two procedures: the forward pass and the backward pass.



Forward Propagation Computation

Input: Data $\{x_i\}$

Output: Loss $\{u_n\}$

```
for  $j = 1, \dots, i$  do
```

```
   $u_j = x_j$ 
```


```
end
```

```
for  $j = i + 1, \dots, n$  do
```

```
   $u_j = f^{(j)}(Pa(u_j))$ 
```

```
end
```

Algorithm 1: Forward Propagation



- Algorithm specifies a computational graph \mathcal{G} .
- Computation in reverse order gives back-propagation computational graph \mathcal{B} .

Backward Propagation Computation

- 1 Proceeds exactly in reverse order of computation in \mathcal{G} .
- 2 Each node in \mathcal{B} computes the derivative $\frac{\partial u_n}{\partial u_i}$.
- 3 Done by using the chain rule w.r.t. the scalar output u_n :

$$\frac{\partial u_n}{\partial u_j} = \sum_{i: j \in Pa(u_i)} \underbrace{\frac{\partial u_n}{\partial u_i}}_{(1)} \cdot \underbrace{\frac{\partial u_i}{\partial u_j}}_{(2)},$$

where all derivatives in the summation can be computed:

- ▶ (1) has been computed in previous step.
- ▶ (2) is easily computed given the mapping from u_j to u_i .
- ▶ acting like dynamic programming.

Backpropagation Algorithm

Input: Data $\{x_i\}$

Output: All derivatives $\{\frac{\partial u_n}{\partial u_j}\}$

Run forward propagation to obtain network activations (outputs)
Initialize grad-table, a data structure that will store derivatives that have been computed, The entry $\text{grad-table}[u_j]$ will store the computed value of $\frac{\partial u_n}{\partial u_j}$

$\text{grad-table}[u_j] \leftarrow 1$

for $j = n - 1, \dots, i + 1$ **do**

$$\quad \text{grad-table}[u_j] = \sum_{i:j \in Pa(u_i)} \text{grad-table}[u_i] \frac{\partial u_i}{\partial u_j} \quad (*)$$

end

Algorithm 2: Backpropagation Algorithm

Step (*) computes $\frac{\partial u_n}{\partial u_j} = \sum_{i:j \in Pa(u_i)} \frac{\partial u_n}{\partial u_i} \cdot \frac{\partial u_i}{\partial u_j}$.

Computational Complexity & Generalization

- 1 Computational cost is proportional to the number of edges in the graph (same as for forward pass):
 - ▶ each edge is visited a limited number of times.
 - ▶ complexity proportional to #edges in the computational graph.
- 2 Backpropagation thus avoids exponential explosion in repeated subexpressions
 - ▶ by simplifications on the computational graph.
- 3 It can be generalized to the case of intermediate tensor-output by replacing the formula $\frac{\partial u_n}{\partial u_j} = \sum_{i:j \in Pa(u_i)} \frac{\partial u_n}{\partial u_i} \cdot \frac{\partial u_i}{\partial u_j}$ in the BP algorithm by the tensor-version:


$$\frac{\partial u_n}{\partial u_j} = \sum_{i:j \in Pa(u_i)} \sum_{\mathbf{k}} (\nabla_{u_j} u_{i\mathbf{k}}) \frac{\partial u_n}{\partial u_{i\mathbf{k}}} ,$$

where \mathbf{k} is the tuple index of a tensor.

Recap: Backward Propagation Computation


$$\frac{\partial u_n}{\partial u_j} = \sum_{i:j \in Pa(u_i)} \underbrace{\frac{\partial u_n}{\partial u_i}}_{(1)} \cdot \underbrace{\frac{\partial u_i}{\partial u_j}}_{(2)},$$

- (1) has been computed in previous step.
- (2) is easily computed given the mapping from u_j to u_i .
- When u_i 's are vectors, (2) becomes a matrix (Jacobian matrix), and (1) becomes a vector: \rightarrow the result is a vector.



Backprop in Fully Connected MLP

- Maps parameters to supervised loss $L(\hat{y}, y)$ associated with a single training example (\mathbf{x}, y) with \hat{y} the output when \mathbf{x} is the input.



Forward Prop for MLP (without regularization)

Input: Network depth l ; Weight matrices $\mathbf{W}^{(i)}, i \in \{1, \dots, l\}$; bias parameters $\mathbf{b}^{(i)}, i \in 1, \dots, l$; Input data \mathbf{x} and target output y ;

Input: Activation function $f^{(k)}, k \in \{1, \dots, l\}$

Output: Loss function L

$$\tilde{\mathbf{h}}^{(0)} = \mathbf{x}$$

for $k = 1, \dots, l$ **do**

$$\begin{array}{|l} \mathbf{a}^{(k)} = \mathbf{W}^{(k)T} \mathbf{h}^{(k-1)} + \mathbf{b}^{(k)} \\ \mathbf{h}^{(k)} = f^{(k)}(\mathbf{a}^{(k)}) \end{array}$$

end

$$\hat{y} = \mathbf{h}^{(l)}$$

Return $L(\bar{y}, y)$

Algorithm 3: Forward prop through typical deep NN

- $L(\bar{y}, y)$ is the loss function w.r.t. the network output \bar{y} and the target y , e.g., the cross entropy.

Backward Prop for MLP

Forward prop through typical deep NN

$\mathbf{g} \leftarrow \nabla_{\bar{y}} \mathcal{J} = \nabla_{\bar{y}} L(\bar{y}, y)$ (gradient of hidden layers, e.g., derivative of softmax on the top-layer)

for $k = l, l-1, \dots, 1$ do

 Calculate the gradient of the pre-nonlinearity activation $\mathbf{a}^{(k)}$ (element-wise multiplication if f is elementwise)

 // \mathbf{g} stores either $\nabla_{\mathbf{a}^{(k)}} \mathcal{J}$ or $\nabla_{\mathbf{h}^{(k)}} \mathcal{J}$

$$\mathbf{g} \leftarrow \nabla_{\mathbf{a}^{(k)}} \mathcal{J} = \nabla_{f(\mathbf{a}^{(k)})} \mathcal{J} \odot \frac{f'(\mathbf{a}^{(k)})}{\mathbf{a}^{(k)}} = \mathbf{g} \odot f'(\mathbf{a}^{(k)})$$

 Compute the gradients on weights and biases (including the regularization term: $\mathbf{a}^{(k)} = \mathbf{W}^{(k)} \mathbf{h}^{(k-1)} + b^{(k)}$)

$$\nabla_{b^{(k)}} \mathcal{J} = \nabla_{\mathbf{a}^{(k)}} \mathcal{J} \odot \frac{\partial \mathbf{a}^{(k)}}{\partial b^{(k)}} = \mathbf{g}$$

$$\nabla_{\mathbf{W}^{(k)}} \mathcal{J} = \nabla_{\mathbf{a}^{(k)}} \mathcal{J} \odot \frac{\partial \mathbf{a}^{(k)}}{\partial \mathbf{W}^{(k)}} = \mathbf{g} \mathbf{h}^{(k-1)T}$$

 Propagate the gradients w.r.t. the next lower-level hidden layer's activations:

$$\mathbf{g} = \nabla_{\mathbf{h}^{(k-1)}} \mathcal{J} = \frac{\partial \mathbf{a}^{(k)}}{\partial \mathbf{h}^{(k-1)}} \cdot \nabla_{\mathbf{a}^{(k)}} \mathcal{J} = \mathbf{W}^{(k)T} \mathbf{g}$$

end

Implementation: Symbol-to-Symbol Derivatives

- 1 Algebraic expressions and computational graphs operate on symbolic representations.
- 2 When we train a neural network we must assign specific values for these symbols, *e.g.*, $[2.5, 3.75, -1.8]^T$.

Two Approaches to BP


1 Symbol-to-number differentiation:

- ▶ Input: computational graph and numerical values of inputs to graph.
- ▶ Output: a set of numerical values describing gradient at the input values.
- ▶ Used by libraries: PyTorch/Torch and Caffe.

2 Symbol-to-symbol differentiation:

- ▶ Input: computational graph.
- ▶ Output: Add additional nodes to the graph.
- ▶ Used by libraries: Theano and Tensorflow.
- ▶ BP does not need to ever access any actual numerical values:
 - ★ instead it adds nodes to a computational graph describing how to compute the derivatives for any specific numerical values.
 - ★ a generic graph evaluation engine can later compute derivatives for any specific numerical values.

Example: Symbol-to-symbol Derivatives



Advantages of Approach

- 1 Derivatives are described in the same language as the original expression.
- 2 Because the derivatives are just another computational graph, it is possible to run back-propagation again:
 - ▶ differentiating the derivatives.
 - ▶ yields higher-order derivatives.

How is BP Implemented

Constructing the Graph

- 1 Each node in graph \mathcal{G} corresponds to a variable.
- 2 Each variable is described by a tensor \mathbf{V} , subsuming scalars, vectors and matrices.
- 3 For each node, several routines are implemented associated with the node.

Subroutines Associated with V

- $\text{get_operation}(V)$:
 - ▶ returns the operation that computes V represented by the edges coming into V .
 - ▶ e.g., suppose we have a variable that is computed by matrix multiplication $C = AB$, then $\text{get_operation}(V)$ returns a pointer to an instance of the corresponding C++ class implemented the multiplication.
- $\text{get_consumers}(V, \mathcal{G})$:
 - ▶ returns list of variables that are children of V in the computational graph \mathcal{G} .
- $\text{get_inputs}(V, \mathcal{G})$:
 - ▶ returns list of variables that are parents of V in the computational graph \mathcal{G} .

bprop Operation

- 1 Each operation (node) *op* is also associated with a *bprop* operation.
- 2 *op.bprop* operation can compute a Jacobian-vector product as described by: $\nabla_{\mathbf{x}}z = \left(\frac{\partial \mathbf{y}}{\partial \mathbf{x}}\right)^T \nabla_{\mathbf{y}}z$ for input \mathbf{x} and output \mathbf{y} .
- 3 This is how the backpropagation algorithm can achieve great generality:
 - ▶ each operation is responsible for knowing how to backpropagate through the edges in the graph that it participates in (local operation).

Inputs, Outputs of bprop

- 1 Backpropagation algorithm itself does not need to know any differentiation rules:
 - ▶ `op.bprop` implements it.
 - ▶ it only needs to call each operation's `op.bprop` rules with the right arguments.
- 2 Formally `op.bprop(inputs, X, D)` must return:

$$\nabla_{\mathbf{x}} z = \left(\frac{\partial \mathbf{y}}{\partial \mathbf{x}} \right)^T \nabla_{\mathbf{y}} z = \sum_i (\nabla_{\mathbf{x}} y_i) \mathbf{D}_i$$

- ▶ “inputs” is a list of inputs that are supplied to the operation, **y** is a math function that the operation implements.
- ▶ **X** is the input whose gradient we wish to compute.
- ▶ **D** is the gradient on the output of the operation, e.g., gradient of last layer.
- ▶ Just an implementation of the chain rule: $\nabla_{\mathbf{x}} z = \left(\frac{\partial \mathbf{y}}{\partial \mathbf{x}} \right)^T \underbrace{\nabla_{\mathbf{y}} z}_{\mathbf{D}_i}$.

Software Implementations

- 1 Usually provide both:
 - ▶ operations.
 - ▶ their *bprop* methods.
- 2 Users of software libraries are able to backpropagate through graphs built using common operations like:
 - ▶ matrix multiplication, exponents, logarithms, *etc.*
- 3 To add a new operation to existing library, one must derive `ob.prop` method manually.

Formal Backpropagation Algorithm: Outermost Skeleton

Require: \mathbb{T} , the target set of variables whose gradients must be computed.

Require: \mathcal{G} , the computational graph

Require: z , the variable to be differentiated

Let \mathcal{G}' be \mathcal{G} pruned to contain only nodes that are ancestors of z and descendants of nodes in \mathbb{T} .

Initialize **grad_table**, a data structure associating tensors to their gradients

grad_table[z] $\leftarrow 1$

for \mathbf{V} in \mathbb{T} **do**

build_grad($\mathbf{V}, \mathcal{G}, \mathcal{G}', \mathbf{grad_table}$)

end for

Return **grad_table** restricted to \mathbb{T}

Inner Loop: build-grad (\mathbf{V} , \mathcal{G} , \mathcal{G}' , grad-table)

Require: \mathbf{V} , the variable whose gradient should be added to \mathcal{G} and grad_table.

Require: \mathcal{G} , the graph to modify.

Require: \mathcal{G}' , the restriction of \mathcal{G} to nodes that participate in the gradient.

Require: grad_table, a data structure mapping nodes to their gradients

if \mathbf{V} is in grad_table then

Return grad_table[\mathbf{V}]

end if

$i \leftarrow 1$

for \mathbf{C} in get_consumers(\mathbf{V} , \mathcal{G}') do

op \leftarrow get_operation(\mathbf{C})

$\mathbf{D} \leftarrow$ build_grad(\mathbf{C} , \mathcal{G} , \mathcal{G}' , grad_table)

$\mathbf{G}^{(i)} \leftarrow$ op.bprop(get_inputs(\mathbf{C} , \mathcal{G}'), \mathbf{V} , \mathbf{D})

$i \leftarrow i + 1$

end for

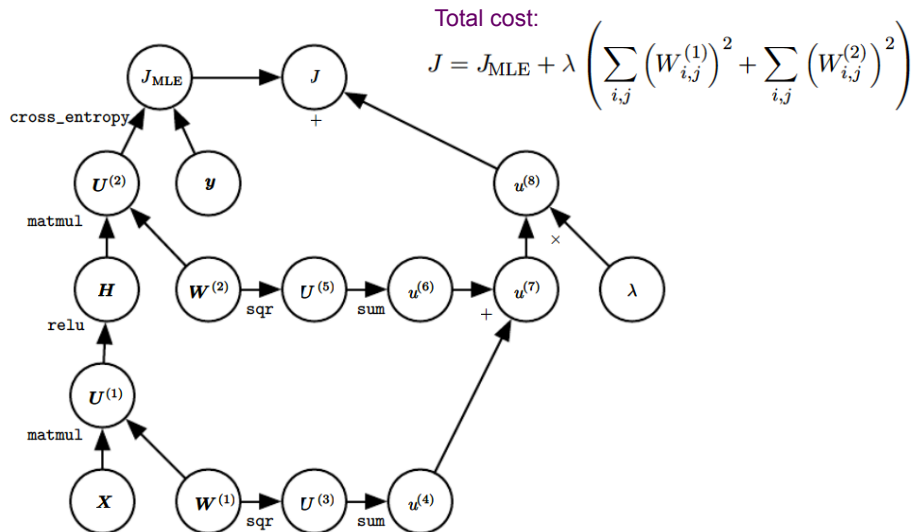
$\mathbf{G} \leftarrow \sum_i \mathbf{G}^{(i)}$

grad_table[\mathbf{V}] = \mathbf{G}

Insert \mathbf{G} and the operations creating it into \mathcal{G}

Return \mathbf{G}

Example: FNN Forward Propagation Graph



Computational Graph of Gradient

- 1 It would be large and tedious for this example.
- 2 One benefit of back-propagation algorithm is that it can automatically generate gradients that would be straightforward but tedious manually for a software engineer to derive.

After gradient computation, it is the responsibility of SGD or other optimization algorithm to use gradients to update parameters.