

# Reinforcement Learning: Basics and DQN

Changyou Chen

Department of Computer Science and Engineering  
University at Buffalo, SUNY  
`changyou@buffalo.edu`


April 25, 2019

# Limitation


## Limitations:

- Iteration over / storage for all states and actions: requires small, discrete state-action space:
  - sampling-based approximations
- Update equations require access to dynamics model:
  - $Q/V$  function fitting.

# Can tabular methods scale?



Gridworld  
 $10^4$



Tetris  
 $10^6$



Atari  
 $10^{308}$  (ram)  $10^{16992}$  (pixels)

## Approximate Q-Learning

- Instead of a table, we have a parametrized Q-function  $Q_\theta(s, a)$ :
  - ▶ Can be a linear function in features:

$$Q_\theta(s, a) = \theta_0 f_0(s, a) + \theta_1 f_1(s, a) + \cdots + \theta_n f_n(s, a)$$

- ▶ Or a complicated neural network.
- Learning rule:

$$\theta_{k+1} = \theta_k - \alpha \nabla_\theta \left[ \frac{1}{2} (Q_\theta(s, a) - \text{target}(s'))^2 \right] \Big|_{\theta=\theta_k}$$

$$\text{target}(s') \triangleq R(s, a, s') + \gamma \max_{a'} Q_{\theta_k}(s', a')$$

$$\begin{aligned} Q_{k+1}(s, a) &= \sum_{s'} P(s'|s, a) \left( R(s, a, s') + \gamma \max_{a'} Q_k(s', a') \right) \\ &\approx R(s, a, s') + \gamma \max_{a'} Q_k(s', a'), \quad \text{with } s' \sim P(\cdot|s, a) \end{aligned}$$

## Connection to Tabular Q-Learning

- Suppose  $\theta \in \mathbb{R}^{|S| \times |A|}$ ,  $Q_\theta(s, a) \equiv \theta_{sa}$

$$\begin{aligned} & \nabla_{\theta_{sa}} \left[ \frac{1}{2} (Q_\theta(s, a) - \text{target}(s'))^2 \right] \\ &= \nabla_{\theta_{sa}} \left[ \frac{1}{2} (\theta_{sa} - \text{target}(s'))^2 \right] \\ &= \theta_{sa} - \text{target}(s') \end{aligned}$$

- Plug into update:  $\theta_{sa} \leftarrow \theta_{sa} - \alpha(\theta_{sa} - \text{target}(s'))$   
 $= (1 - \alpha)\theta_{sa} + \alpha[\text{target}(s')]$

- Compare with Tabular Q-Learning update:

$$Q_{k+1}(s, a) \leftarrow (1 - \alpha)Q_k(s, a) + \alpha[\text{target}(s')]$$

## Convergence of Approximate Q-Learning

- It is not guaranteed to converge, even if the function approximation is expressive enough to represent the true  $Q$ -function:
  - ▶ The approximation is sequential, so if each time there induces a large enough error, the aggregate error might be exploded.


# Deep Q-Learning<sup>1</sup>

<sup>1</sup>Adapted from [http://www.teach.cs.toronto.edu/~csc2542h/fall/material/csc2542f16\\_dqn.pdf](http://www.teach.cs.toronto.edu/~csc2542h/fall/material/csc2542f16_dqn.pdf)

# Artari Games



Pong



Breakout



Space Invaders



Seaquest



Beam Rider

## Breakout game

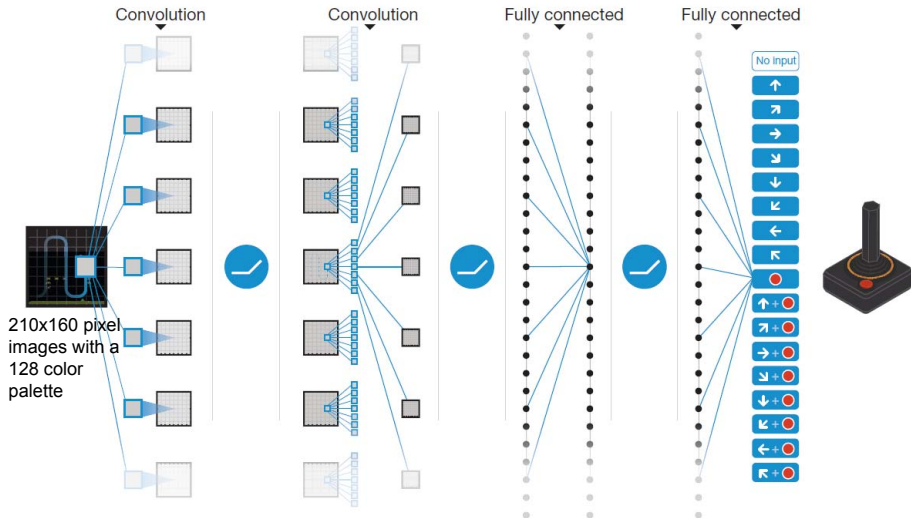
How to define a state?

- Location of the paddle
- Location/direction of the ball
- Presence/absence of each individual brick

Use screen pixels!



# DQN in Atari



- Naive Q-learning oscillates or diverges with neural nets:
  - ▶ Data is sequential:
    - Successive samples are correlated, non-i.i.d.
  - ▶ Policy changes rapidly with slight changes to Q-values:
    - Policy may oscillate
    - Distribution of data can swing from one extreme to another
  - ▶ Scale of rewards and Q-values is unknown:
    - Naive Q-learning gradients can be large unstable when backpropagated

## Approximate Q-Learning

- Instead of a table, we have a parametrized Q-function  $Q_\theta(s, a)$ :
  - ▶ Can be a linear function in features:

$$Q_\theta(s, a) = \theta_0 f_0(s, a) + \theta_1 f_1(s, a) + \cdots + \theta_n f_n(s, a)$$

- ▶ Or a complicated neural network.
- Learning rule:

$$\theta_{k+1} = \theta_k - \alpha \nabla_\theta \left[ \frac{1}{2} (Q_\theta(s, a) - \text{target}(s'))^2 \right] \Big|_{\theta=\theta_k}$$

$$\text{target}(s') \triangleq R(s, a, s') + \gamma \max_{a'} Q_{\theta_k}(s', a')$$

## Deep Q-Learning

- Deep Q-Network provides a stable solution to deep value-based RL:

- ▶ Use experience replay:
  - Break correlations in data, bring us back to i.i.d. setting
  - Learn from all past policies
  - Using off-policy Q-learning
- ▶ Freeze target Q-network:
  - Avoid oscillations
  - Break correlations between Q-network and target
- ▶ Clip rewards or normalize network adaptively to sensible range:
  - Robust gradients

Algorithm:

Start with  $Q_0(s, a)$  for all  $s, a$ .

Get initial state  $s$

For  $k = 1, 2, \dots$  till convergence

    Sample action  $a$ , get next state  $s'$

    If  $s'$  is terminal:

$\text{target} = R(s, a, s')$

        Sample new initial state  $s'$

    else:

$\text{target} = R(s, a, s') + \gamma \max_{a'} Q_k(s', a')$

$Q_{k+1}(s, a) \leftarrow (1 - \alpha)Q_k(s, a) + \alpha [\text{target}]$

$s \leftarrow s'$

# Deep Q-Learning: Experience Replay

- To remove correlations, build data-set from agent's own experience:
  - ▶ Take action  $a_t$  according to  $\epsilon$ -greedy policy
  - ▶ Store transition  $(s_t, a_t, r_{t+1}, s_{t+1})$  in replay memory  $\mathcal{D}$  (Huge data base to store historical samples)
  - ▶ Sample random mini-batch of transitions  $(s, a, r, s')$  from  $\mathcal{D}$
  - ▶ Optimize MSE between  $Q$ -network and  $Q$ -learning targets, e.g.,

$$L_k(\theta_k) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}} \left[ \frac{1}{2} (Q_{\theta_k}(s, a) - \text{target}(s'))^2 \right]$$

$$\text{target}(s') \triangleq R(s, a, s') + \gamma \max_{a'} Q_{\theta_k}(s', a')$$

# Deep Q-Learning: Fixed Target Q-Network

- To avoid oscillations, fix parameters used in Q-learning target:
  - ▶ Compute Q-learning targets w.r.t. old, fixed parameters  $\theta_k^-$ :

$$\text{target}(s'; \theta_k^-) \triangleq R(s, a, s') + \gamma \max_{a'} Q_{\theta_k^-}(s', a')$$

- ▶ Optimize MSE between Q-network and Q-learning targets, e.g.,

$$L_k(\theta_k) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}} \left[ \frac{1}{2} (Q_{\theta_k}(s, a) - \text{target}(s'; \theta_k^-))^2 \right]$$

- ▶ Periodically update fixed parameters  $\theta_k^- = \theta_k$ .

## Deep Q-Learning: Reward / Value Range

- DQN clips the reward to  $[-1, +1]$ .
- This prevents Q-values from becoming too large.
- Ensures gradients are well-conditioned.

## DQN

Game	With replay, with target Q	With replay, without target Q	Without replay, with target Q	Without replay, without target Q
Breakout	316.8	240.7	10.2	3.2
Enduro	1006.3	831.4	141.9	29.1
River Raid	7446.6	4102.8	2867.7	1453.0
Seaquest	2894.4	822.6	1003.0	275.8
Space Invaders	1088.9	826.3	373.2	302.0



# Algorithm

## Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory  $D$  to capacity  $N$

Initialize action-value function  $Q$  with random weights  $\theta$

Initialize target action-value function  $\hat{Q}$  with weights  $\theta^- = \theta$

**For** episode = 1,  $M$  **do**

    Initialize sequence  $s_1 = \{x_1\}$  and preprocessed sequence  $\phi_1 = \phi(s_1)$

**For**  $t = 1, T$  **do**

        With probability  $\varepsilon$  select a random action  $a_t$

        otherwise select  $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

        Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$

        Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$

        Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $D$

        Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $D$

        Set  $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

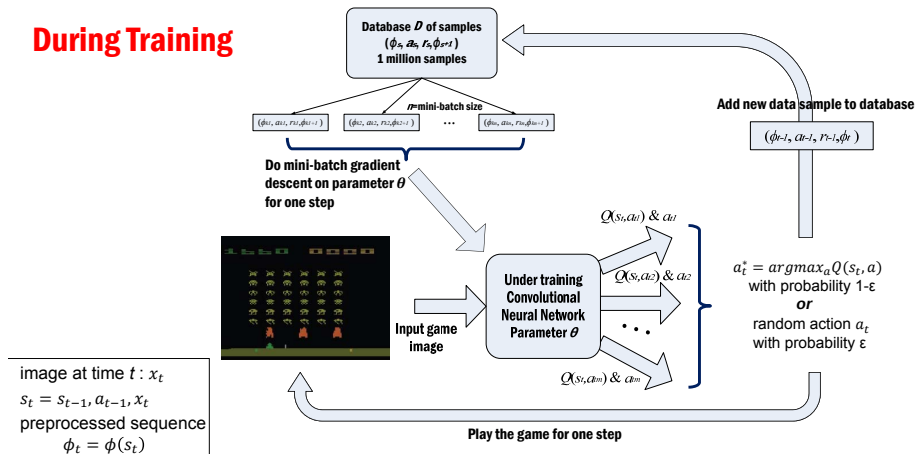
        Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  with respect to the network parameters  $\theta$

        Every  $C$  steps reset  $\hat{Q} = Q$

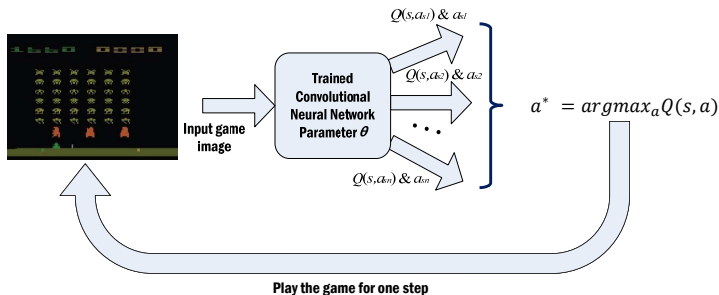
**End For**

**End For**

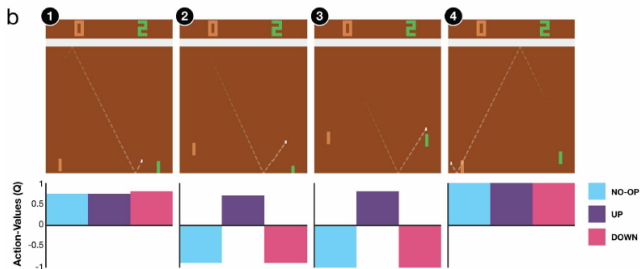
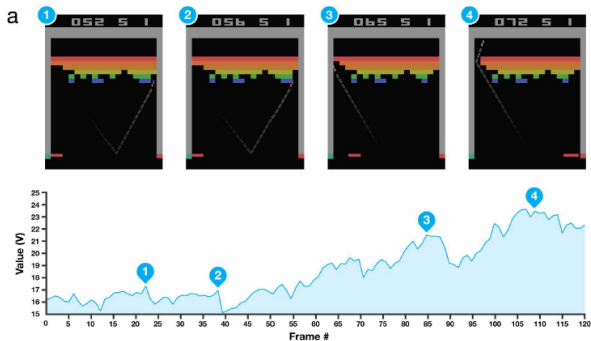
## During Training




## After Training



# Results



# Results



# Double Q-Learning

## Double Q-Learning

- Train 2 action-value functions,  $Q_1$  and  $Q_2$ .
- Do Q-learning on both, but
  - ▶ never on the same time steps ( $Q_1$  and  $Q_2$  are independent)
  - ▶ pick  $Q_1$  or  $Q_2$  at random to be updated on each step
- If updating  $Q_1$ , use  $Q_2$  to evaluate the target:

$$Q_1(s_t, a_t) = Q_1(s_t, a_t) + \alpha \left( R_{t+1} + Q_2(s_{t+1}, \underset{a}{\operatorname{argmax}} Q_1(s_{t+1}, a)) - Q_1(s_t, a_t) \right)$$

- Action selections are  $\epsilon$ -greedy with respect to the sum of  $Q_1$  and  $Q_2$ .

## Deep Q-Learning

$$\text{target}(s'; \theta_k^-) \triangleq R(s, a, s') + \gamma \max_{a'} Q_{\theta_k^-}(s', a')$$

$$L_k(\theta_k) = \mathbb{E}_{s, a, r, s' \sim \mathcal{D}} \left[ \frac{1}{2} (Q_{\theta_k}(s, a) - \text{target}(s'; \theta_k^-))^2 \right]$$

## Double Q-Learning

Initialize  $Q_1(s, a)$  and  $Q_2(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily

Initialize  $Q_1(\text{terminal-state}, \cdot) = Q_2(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

    Initialize  $S$

    Repeat (for each step of episode):

        Choose  $A$  from  $S$  using policy derived from  $Q_1$  and  $Q_2$  (e.g.,  $\varepsilon$ -greedy in  $Q_1 + Q_2$ )

        Take action  $A$ , observe  $R, S'$

        With 0.5 probability:

$$Q_1(S, A) \leftarrow Q_1(S, A) + \alpha \left( R + \gamma Q_2(S', \arg \max_a Q_1(S', a)) - Q_1(S, A) \right)$$

    else:

$$Q_2(S, A) \leftarrow Q_2(S, A) + \alpha \left( R + \gamma Q_1(S', \arg \max_a Q_2(S', a)) - Q_2(S, A) \right)$$

$S \leftarrow S'$ ;

    until  $S$  is terminal



## Double Q-Learning


- Current Q-network  $\theta$  is used to select actions.
- Older Q-network  $\theta^-$  is used to evaluate actions.

Action evaluation:  $\theta^-$

$$L = \left( r + \gamma \underbrace{Q(s', \underbrace{\operatorname{argmax}_{a'} Q(s', a', \mathbf{w})}_{\text{Action selection: } \theta})}_{\text{Action evaluation: } \theta^-}, \theta^- \right) - Q(s, a, \theta) \Big)^2$$

Action selection:  $\theta$

# Double Q-Learning



# Dueling Network Architecture

## **$Q$ function should be designed more wisely: containing an action-independent component**

For many states:

- unnecessary to estimate the value of each action choice, for example, move left or right only matters when a collision is eminent.
- In most of states, the choice of action has no affect on what happens

## Decompose $Q$

$$Q^\pi(s, a) = V^\pi(s) + A^\pi(s, a)$$


$$\Rightarrow A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s) = Q^\pi(s, a) - \mathbb{E}_{a \sim \pi}[Q^\pi(s, a)]$$

(advantage function)


# Saliency map on the Atari game Enduro

1. Focus on **horizon**,  
where new **cars** appear

2. Focus on the **score**



**Not pay much attention**  
when there are **no cars** in front




Attention on car immediately in front  
making **its choice of action** very relevant

# Dueling Network

- Single Q-network with two streams.
- Produce separate estimations of state value function  $V$  and advantage function  $A$ .

sharing convolutional feature learning module



- outputs a scalar value  $V(s)$  and a  $|\mathcal{A}|$ -dimensional vector  $A(s, a)$ .

## Aggregation Modules

- Add:  $Q(s, a) = V(s) + A(s, a)$ .
- Subtract max:  $Q(s, a) = V(s) + (A(s, a) - \max_{a'} A(s, a'))$ :
  - ▶ force  $A$  to have zero at the chosen optimal action.
  - ▶ Equivalent to shifting  $V$ .
- Subtract mean:  $Q(s, a) = V(s) + \left( A(s, a) - \frac{1}{|\mathcal{A}|} \sum_{a'} A(s, a') \right)$ :
  - ▶ loses the original semantics of  $V$  and  $A$ .
  - ▶ but increases the stability of the optimization.
  - ▶ often works best.



*Table 1.* Mean and median scores across all 57 Atari games, measured in percentages of human performance.

	<b>30 no-ops</b>		<b>Human Starts</b>	
	<b>Mean</b>	<b>Median</b>	<b>Mean</b>	<b>Median</b>
Prior. Duel Clip	<b>591.9%</b>	<b>172.1%</b>	<b>567.0%</b>	<b>115.3%</b>
Prior. Single	434.6%	123.7%	386.7%	112.9%
Duel Clip	<b>373.1%</b>	<b>151.5%</b>	<b>343.8%</b>	<b>117.1%</b>
Single Clip	341.2%	132.6%	302.8%	114.1%
Single	307.3%	117.8%	332.9%	110.9%
Nature DQN	227.9%	79.1%	219.6%	68.5%